

Devang Borkar

+1 (312) 358-5722 | devangborkar3@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Skills

- **Languages:** Python, C#, Java, TypeScript/JavaScript, SQL, Go, C++
- **AI/ML:** LLMs, Multi-agent systems, LoRA fine-tuning, Prompt engineering, ReACT agents, tool-calling, RAG, embeddings, Harness Engineering, Synthetic Data Generation, Agentic Memory, FAISS
- **Frameworks:** PyTorch, HuggingFace Transformers, FastAPI, LangChain, React, MongoDB, Redis
- **Cloud/APIs:** CI/CD Pipelines, GCP, AWS, Docker, OpenAI/Anthropic/ Gemini APIs, OpenRouter
- **Coding Agents:** Cursor IDE, Claude Code, Codex CLI, OpenCode, Github Copilot

Education

- M.S. Computer Science, University of California, Davis (GPA: 3.91) *Expected May 2026*
- B.E. Computer Science, SPPU, Pune (GPA: 4.0) *Graduated*

Experience

AI Engineer *PilotCrew AI*

Oct 2025 – Present

- Developed end-to-end AI evaluation infrastructure for agentic systems, using Python APIs, MongoDB persistence, Dockerized execution, and observability workflows for repeatable model-quality measurement.
- Built an automated prompt optimization pipeline that analyzes agentic traces and iteratively improve prompts, achieving a 6.63% performance lift for conversational agents on MultiChallenge benchmark.
- Engineered agent harnesses for execution, tool-use tracing, hidden-state management, deterministic grading, and failure analysis across coding, interactive tool-calling and deep research workflows.

SWE Intern *LearnHaus AI*

Jun 2025 – Aug 2025

- Built a full-stack multimodal AI application (React, Python) from scratch, processing video, audio, and text inputs with optimized latency for real-time user feedback.
- Developed a LLM-as-Judge evaluation framework using GPT-5 and Gemini-3 to score model outputs across multiple quality dimensions; deployed to Google Cloud Platform(GCP)

Software Engineer *Hexaview Technologies*

Aug 2022 – Sep 2024

- Shipped 20+ production features for a Fortune 500 fintech platform; contributed to scalable microservices using C# and SQL. Collaborated within a 30+ person cross-functional Agile team using JIRA.
- Optimized a legacy .NET backend servicing over 1 million monthly requests, applying key design patterns to reduce code complexity & successfully redesigning 50+ REST APIs in MVC architecture

Projects

CausalFlow: Agentic Trace Debugging Framework

[GitHub](#) | *NeurIPS '26 submission*

- Built a Python framework for causal attribution in multi-step LLM agent traces, identifying failure-inducing reasoning/tool steps and generating localized repair pairs that are useful as DPO post-training data.
- Evaluated across 3,000+ math, code generation, web QA, and medical browsing tasks; repaired 42.7% of failed agent outputs and improved post-repair accuracy by a median +16.9 % over baseline.

Process Reward Model for On-Device LLMs

[GitHub](#) | *Graduate Research*

- Performed Supervised Fine Tuning on an 8B parameter LLM using PyTorch to verify reasoning steps; optimized for edge deployment with a lightweight 0.6B model.
- Improved model accuracy by 21% on GSM8K using speculative decoding over baseline methods; published model weights to HuggingFace.

ResShare: Decentralized File Sharing Platform

[GitHub](#) | *Apache Foundation - ResilientDB*

- Built a full-stack file-sharing platform using Flask and React, implementing REST APIs for user authentication, access control, file upload/download, folder hierarchy management and sharing workflows.
- Added a document indexing and RAG pipeline for files using text extraction, semantic chunking, embeddings, FAISS search, source-attributed responses, and per-user retrieval isolation.