

Devang Borkar

Davis, CA | +1 (312) 358-5722 | devangborkar3@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Skills

- **Languages:** Python, Go, TypeScript/JavaScript, C#, Java, SQL, C++
- **Backend/Infra:** REST APIs, FastAPI, Flask, PostgreSQL, MongoDB, Redis, Docker, Observability, distributed systems
- **AI/ML Systems:** LLM evaluation, agentic workflows, tool-calling, RAG, embeddings/vector search, deterministic graders, prompt optimization, LLM-as-Judge
- **Frameworks:** PyTorch, HuggingFace Transformers, FastAPI, LangChain, React, MongoDB, Redis
- **Cloud/APIs:** AWS, GCP, OpenRouter, OpenAI/Anthropic/Gemini APIs, GitHub Actions, CI/CD
- **Coding Agents:** Claude Code, Codex CLI, Cursor IDE, OpenCode, Github Copilot

Education

University of California, Davis

Masters of Science in Computer Science

Sept 2024 - June 2026

GPA: 3.91/4.0

Relevant Coursework: *Distributed Databases, Machine Learning, Data Structures and Algorithms, Vision Language Models*

Pune University, India

Bachelor in Computer Engineering | Honors Track in AI/ML

Aug 2019 - June 2023

GPA: 4.0/4.0

Experience

Software Engineer

PilotCrew AI

Oct 2025 – Present

Davis, CA

- Developed end-to-end AI evaluation infrastructure for agentic systems, using Python APIs, MongoDB, Docker, and observability workflows for repeatable model-quality measurement.
- Built a prompt optimization pipeline that analyzes agentic traces and iteratively improve prompts, achieving a 6.63% performance lift for conversational agents on MultiChallenge benchmark.
- Engineered agent harnesses for execution, tool-use tracing, hidden-state management, deterministic grading, and failure analysis across coding, interactive tool-calling and deep research workflows.

Software Engineer Intern

LearnHaus AI

Jun – Aug 2025

Sacramento, CA

- Built a full-stack multimodal AI coaching application in React and Python, processing video, audio, and transcript signals with async pipelines for real-time user feedback.
- Developed an LLM-as-Judge evaluation framework using GPT-5 and Gemini-3 to score response quality across multiple dimensions, deploying the service on GCP.

Software Engineer

Hexaview Technologies

Aug 2022 – Sept 2024

Pune, India

- Shipped 50+ production features for a Fortune 500 fintech platform; contributed to scalable microservices using C# and SQL. Collaborated within a 30+ person cross-functional Agile team using JIRA.
- Optimized a legacy .NET backend servicing over 1 million monthly requests, applying key design patterns to reduce code complexity & successfully redesigning 50+ REST APIs in MVC architecture

Projects

Process Reward Model for On-Device LLMs

[GitHub](#) | Graduate Research

- Performed Supervised Fine Tuning on an 8B parameter LLM using PyTorch to verify reasoning steps; optimized for edge deployment with a lightweight 0.6B model.
- Improved model accuracy by 21% on GSM8K using speculative decoding over baseline methods; published model weights to HuggingFace.

ResShare: Decentralized File Sharing Platform

[GitHub](#) | Apache Foundation - ResilientDB

- Built a full-stack file-sharing platform using Flask and React, implementing REST APIs for user authentication, access control, file upload/download, folder hierarchy management and sharing workflows.
- Added a document indexing and RAG pipeline for files using text extraction, semantic chunking, embeddings, FAISS search, source-attributed responses, and per-user retrieval isolation.

CausalFlow: Agentic Trace Debugging Framework

| [GitHub](#) | NeurIPS '26 submission

- Built a Python framework for causal attribution in multi-step LLM agent traces, identifying failure-inducing reasoning/tool steps and generating localized repair pairs that are useful as DPO post-training data.
- Evaluated across 3,000+ math, code generation, web QA, and medical browsing tasks; repaired 42.7% of failed agent outputs and improved post-repair accuracy by a median +16.9 % over baseline.